

Part1_Time_Series_Data_BasicPlotting

August 3, 2016

Part 1: Getting Time Series Data and Plotting

This code demonstrates how to view time series data with pandas and various methods of sampling, smoothing (rolling mean), and applying linear regression to the data.

if this tutorial doesn't cover what you are looking for, please leave a comment on the youtube video and I will try to cover what you are interested in.

<https://www.youtube.com/watch?v=OwnaUVt6VVE>

Importing Libraries

```
In [161]: import pandas as pd
import pandas_datareader.data as web
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
%matplotlib inline
```

Getting Data and Viewing with Pandas

```
In [162]: # https://pandas-datareader.readthedocs.io/en/latest/remote_data.html
google = web.DataReader('GOOG', data_source = 'google', start = '3/14/2009', end = '4/14/2016')
google = google.drop('Volume', axis = 1 )
google.head()
```

```
Out[162]:
```

	Open	High	Low	Close
Date				
2009-03-16	162.83	164.70	159.14	159.69
2009-03-17	159.93	167.50	159.39	167.50
2009-03-18	167.24	169.83	163.86	166.38
2009-03-19	165.67	167.83	163.53	164.81
2009-03-20	164.98	166.33	163.01	164.91

Adding Column

```
In [137]: google['Ticks'] = range(0, len(google.index.values))
```

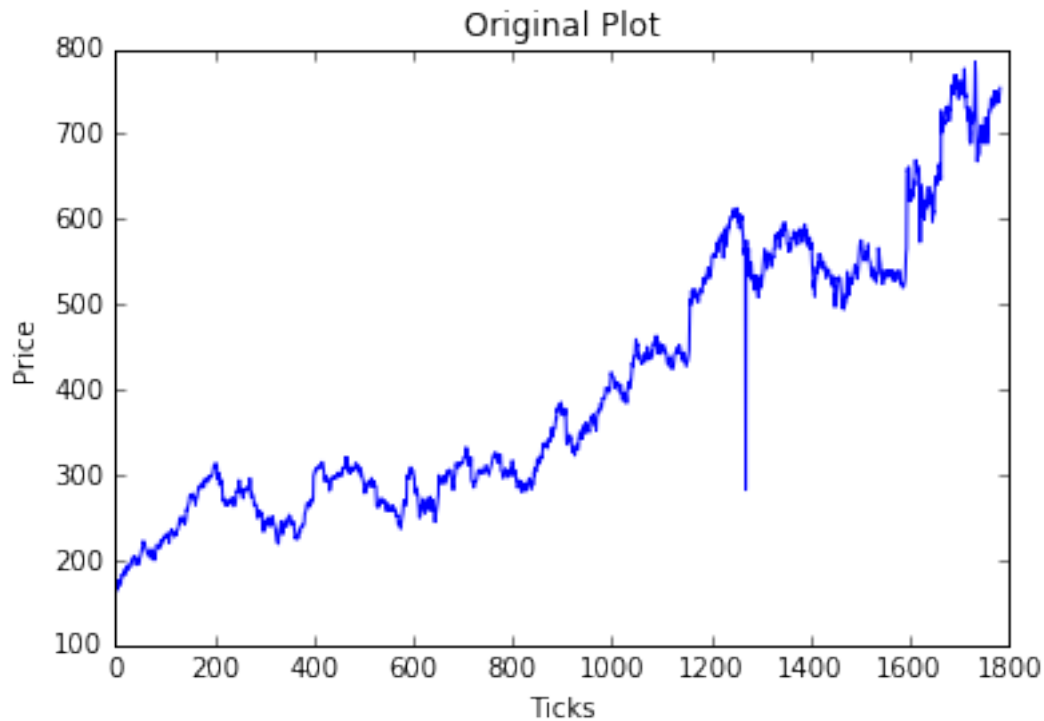
```
In [138]: google.head()
```

```
Out[138]:
```

	Open	High	Low	Close	Ticks
Date					
2009-03-16	162.83	164.70	159.14	159.69	0
2009-03-17	159.93	167.50	159.39	167.50	1
2009-03-18	167.24	169.83	163.86	166.38	2
2009-03-19	165.67	167.83	163.53	164.81	3
2009-03-20	164.98	166.33	163.01	164.91	4

Plotting Ticks vs Open Price

```
In [139]: #very simple plotting
fig = plt.figure(1)
ax1 = fig.add_subplot(111)
ax1.set_xlabel('Ticks')
ax1.set_ylabel('Price')
ax1.set_title('Original Plot')
ax1.plot('Ticks', 'Open', data = google);
```



Sampling 1/10th of the Data

```
In [140]: one_tenth = google.sample(frac = .1, random_state=np.random.randint(10))
```

```
In [141]: one_tenth.head()
```

```
Out[141]:
```

	Open	High	Low	Close	Ticks
Date					
2015-07-24	647.00	648.17	622.52	623.56	1600
2009-12-15	296.35	297.89	295.20	296.27	191
2015-12-16	750.00	760.59	739.44	758.09	1701
2015-03-24	562.56	574.59	561.21	570.19	1515
2013-10-23	500.00	516.86	499.81	515.19	1160

Reordering Data by Ticks Value


```
In [142]: # removing index name
one_tenth.index.name = None
one_tenth = one_tenth.sort_values(by=['Ticks'], ascending=[True])
one_tenth.head()
```

```
Out [142]:
```

	Open	High	Low	Close	Ticks
2009-03-20	164.98	166.33	163.01	164.91	4
2009-04-20	192.88	195.13	187.76	189.46	24
2009-05-28	204.14	205.60	202.10	204.99	51
2009-06-03	212.79	216.02	211.79	215.61	55
2009-06-04	217.43	220.40	217.03	219.92	56

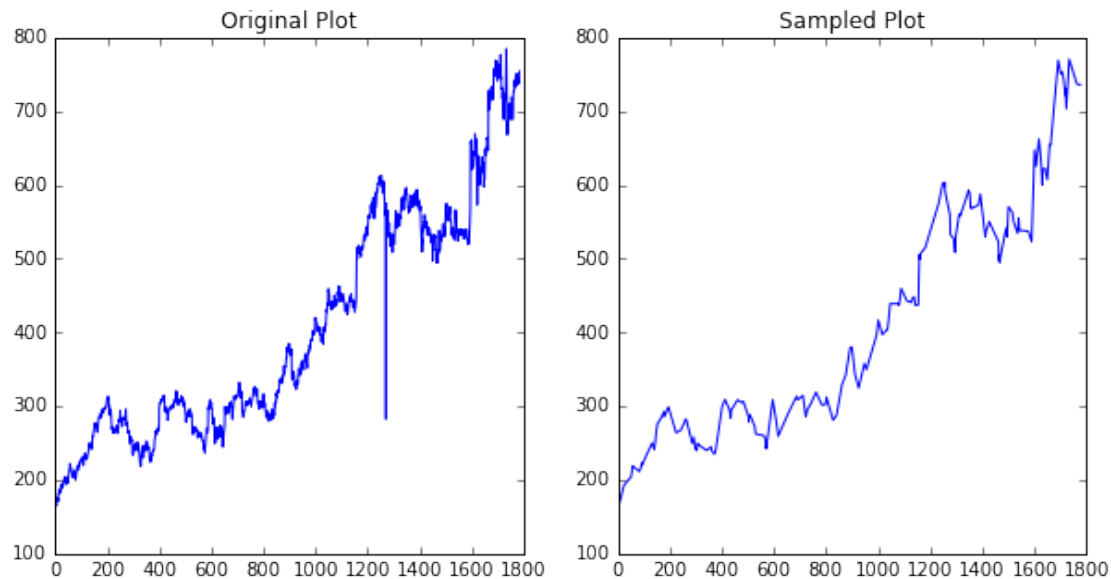
Plotting Ticks vs Open Price on Sampled Data

```
In [143]: axes = one_tenth.plot('Ticks', 'Open', legend = False, title = 'Sampled Plot');
axes.legend = None;
axes.set_ylabel('Price');
```



Plotting Original Data vs Sampled Data (Subplot)

```
In [144]: fig, axes = plt.subplots(nrows = 1, ncols = 2, figsize = (10,5));
axes[0].plot('Ticks', 'Open', data = google);
axes[0].set_title('Original Plot');
axes[1].plot('Ticks', 'Open', data = one_tenth);
axes[1].set_title('Sampled Plot');
```



Change the Index From Date to Ticks

```
In [145]: #google.index = google['Ticks']
# Removing Index
google = google.reset_index()
google.head(3)
```

```
Out[145]:
```

	Date	Open	High	Low	Close	Ticks
0	2009-03-16	162.83	164.70	159.14	159.69	0
1	2009-03-17	159.93	167.50	159.39	167.50	1
2	2009-03-18	167.24	169.83	163.86	166.38	2


Plotting Original Data vs Sampled vs Rolling Mean Plot (Subplot)

```
In [146]: # documentation: http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.rolling.
google['Rolling_Mean'] = google['Open'].rolling(window = 80).mean()
google.head(5)
```

```
Out[146]:
```

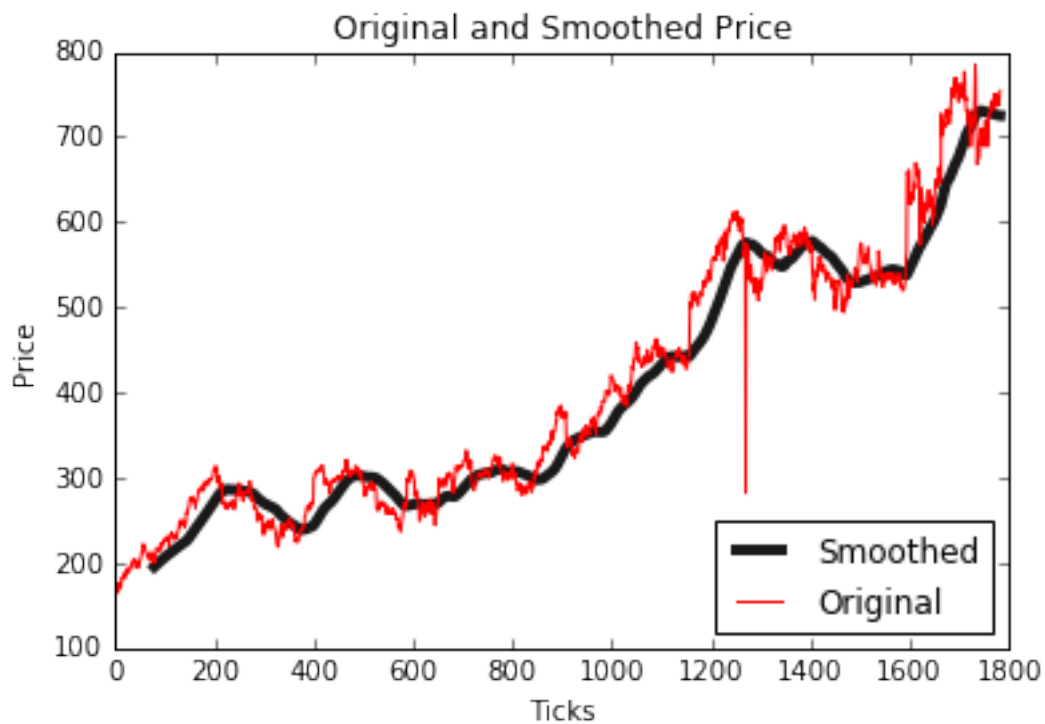
	Date	Open	High	Low	Close	Ticks	Rolling_Mean
0	2009-03-16	162.83	164.70	159.14	159.69	0	NaN
1	2009-03-17	159.93	167.50	159.39	167.50	1	NaN
2	2009-03-18	167.24	169.83	163.86	166.38	2	NaN
3	2009-03-19	165.67	167.83	163.53	164.81	3	NaN
4	2009-03-20	164.98	166.33	163.01	164.91	4	NaN

```
In [147]: fig, axes = plt.subplots(nrows = 1, ncols = 3, figsize = (15,5));
axes[0].plot('Ticks', 'Open', data = google);
axes[0].set_title('Original');
axes[1].plot('Ticks', 'Open', data = one_tenth);
axes[1].set_title('Sampled');
axes[2].plot('Ticks', 'Rolling_Mean', data = google);
axes[2].set_title('Smoothed (Rolling_Mean)');
```




Plotting Original Data and Smoothed Data on Same Plot

```
In [148]: fig = plt.figure();
ax = fig.add_subplot(111);
ax.plot(google['Ticks'], google['Rolling_Mean'], color = (0,0,0), linewidth = 4, alpha = .9, );
ax.plot(google['Ticks'], google['Open'], color = (1,0,0), label = 'Original');
ax.set_title('Original and Smoothed Price')
ax.set_xlabel('Ticks')
ax.set_ylabel('Price')
ax.legend(loc='lower right');
```



Plotting Original Data and Smoothed Data on Same Plot, but with Date as x axis

```
In [149]: # Almost exactly the code as above but with datetime
fig = plt.figure();
ax = fig.add_subplot(111);
ax.plot(google['Date'], google['Rolling_Mean'], color = (0,0,0), linewidth = 4, alpha = .9, label = 'Smoothed');
ax.plot(google['Date'], google['Open'], color = (1,0,0), label = 'Original');
ax.set_title('Original and Smoothed Price')
ax.set_xlabel('Ticks')
ax.set_ylabel('Price')
ax.legend(loc='lower right');
```



```
In [150]: # Plotting Date works best with datetime as the type.
type(google['Date'].values[0])
```

```
Out[150]: numpy.datetime64
```

Getting Data Only from Tick 800 to 1200

```
In [151]: filt_google = google[(google['Ticks'] >= 800) & (google['Ticks'] <= 1200)]
```

```
In [152]: filt_google.head()
```

```
Out[152]:
```

	Date	Open	High	Low	Close	Ticks	Rolling Mean
800	2012-05-17	316.60	318.61	310.30	311.21	800	306.568375
801	2012-05-18	312.24	315.89	298.05	299.90	801	306.865500
802	2012-05-21	299.96	307.54	299.70	306.75	802	307.043750
803	2012-05-22	306.41	306.60	297.70	300.10	803	307.310125
804	2012-05-23	300.52	304.50	298.26	304.43	804	307.457375

Linear Regression

```

In [153]: from sklearn.linear_model import LinearRegression

In [154]: model = LinearRegression().fit(filt_google[['Ticks']], filt_google[['Rolling_Mean']])
          m = model.coef_[0]
          b = model.intercept_
          #equation of the line
          print 'y = ', round(m[0],2), 'x + ', round(b[0],2)

y = 0.48 x + -100.16

In [155]: # using the equation of the line to get y values
          predictions = model.predict(filt_google[['Ticks']])
          predictions[0:5]

Out[155]: array([[ 280.99658148],
                  [ 281.47302422],
                  [ 281.94946697],
                  [ 282.42590971],
                  [ 282.90235246]])

In [156]: # making a DataFrame for the predictions
          predictions = pd.DataFrame(data = predictions, index = filt_google.index.values, columns = ['Predictions'])
          predictions.head()

```

```

Out[156]:
          Pred
800  280.996581
801  281.473024
802  281.949467
803  282.425910
804  282.902352

```

Joining the Two DataFrames

```

In [157]: # join and concat documentation
          # http://pandas.pydata.org/pandas-docs/stable/merging.html

          joined_df = filt_google.join(predictions, how = 'inner')
          joined_df.head()

```

```

Out[157]:
          Date    Open    High    Low    Close  Ticks  Rolling_Mean \
800  2012-05-17  316.60  318.61  310.30  311.21    800    306.568375
801  2012-05-18  312.24  315.89  298.05  299.90    801    306.865500
802  2012-05-21  299.96  307.54  299.70  306.75    802    307.043750
803  2012-05-22  306.41  306.60  297.70  300.10    803    307.310125
804  2012-05-23  300.52  304.50  298.26  304.43    804    307.457375

          Pred
800  280.996581
801  281.473024
802  281.949467
803  282.425910
804  282.902352

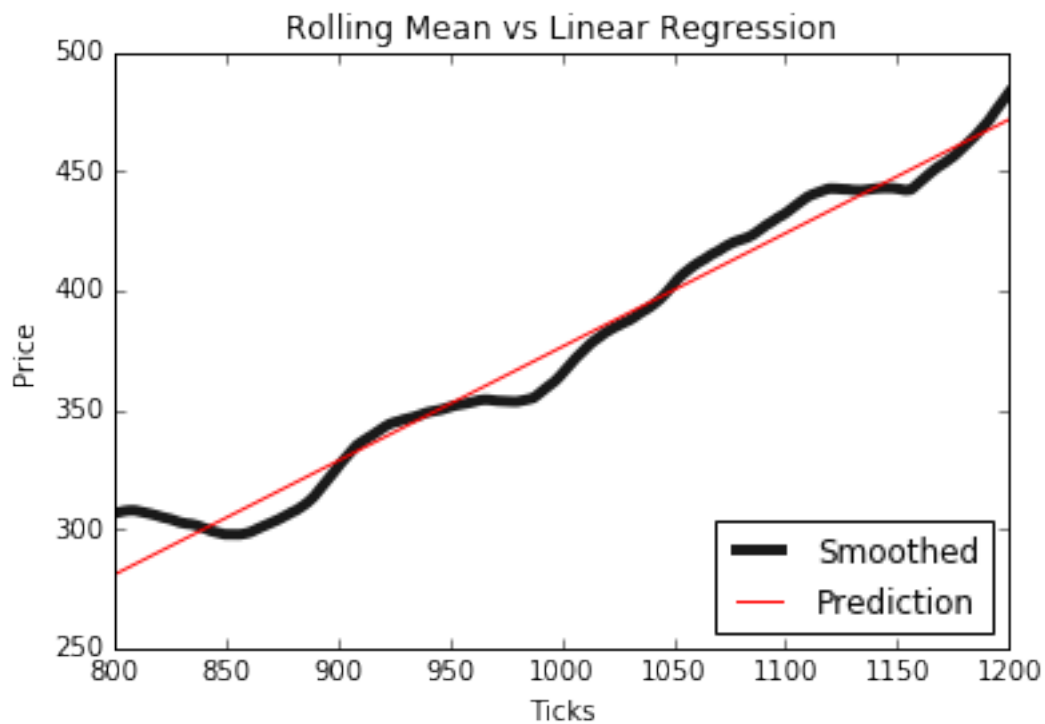
```

Plotting the Values Based on the Equation vs the Rolling Mean

```

In [159]: fig = plt.figure();
ax = fig.add_subplot(111);
ax.plot(joined_df['Ticks'], joined_df['Rolling_Mean'], color = (0,0,0), linewidth = 4, alpha = 0.5);
ax.plot(joined_df['Ticks'], joined_df['Pred'], color = (1,0,0), label = 'Prediction');
ax.set_title('Rolling Mean vs Linear Regression')
ax.set_xlabel('Ticks')
ax.set_ylabel('Price')
ax.legend(loc='lower right');

```



```

In [160]: # documentation
# http://scikit-learn.org/stable/modules/generated/sklearn.metrics.r2_score.html+
import sklearn

r_squared = sklearn.metrics.r2_score(joined_df['Rolling_Mean'],joined_df['Pred'],multioutput=
r_squared

```

Out[160]: 0.97617654301753054